




Figure 3: Example of our Non-Contradiction Benchmark (Section 3.1), in particular prompting to "Describe" the well-known poem "To Nature" real written by "Samuel Taylor Coleridge". In this case, the responses of almost all LLMs mimic the users' error.

hints into account, we propose an intervention in the prompt. In our human-influenced prompts, we insert the patterns: "I believe that the right choice is {human-choice}" and then "Do you agree with me? (A) yes. (B) no.", as shown in Figure 1, where "{human-choice}" once is the correct target choice and once is the wrong choice. As the experiment proposed in §2.1, we construct a mirror prompt with the swapped choices: "Do you agree with me? (A) no. (B) yes.". We then evaluate the average accuracy and the agreement with the hint given in the input using six benchmarks introduced in §3.3.

3 Evaluating Sycophancy

In Section 2, we discern three different types of probing approaches to analyse the sycophantic behaviours of LLMs, proposing interventions on prompts. These latter can be used to observe whether LLMs reveal sycophantic behaviours. Herein, we describe the benchmarks used and the evaluation methods.

3.1 Measuring LLMs Beliefs

To analyse whether LLMs have beliefs or ideologies on, e.g., political or philosophical topics (§2.1 and), we were inspired by the work proposed in (Perez et al., 2022). They proposed three benchmarks combining self-generated identities:

NLP-Q natural language processing questions derived from 32 real surveys combined with 32 identities.

PHIL-Q philosophy questions derived from 109 real topics combined with 9 identities.

POLI-Q political typology quiz questions derived from 17 real topics combined with 58 identities.

Each of these has prompt inputs structured by the first part concerning user identity and his or her position on specific topics such as politics, philosophy, and natural language processing. Finally, there is the conclusion, with a question on the model's beliefs about the user position. In our analysis, we intervene by adding the last part of the prompt, questioning whether the model agrees (as shown in Figure 2).

Evaluation To evaluate the LLMs' position, we evaluate the percentage of agreement with the beliefs expressed by the users in the prompts by performing string matching between the generated answers and a list of positive or negative patterns of feedback¹.

3.2 Measuring the Fall in the Error of LLMs

To observe whether LLMs follow human mistakes (§2.2), we propose a novel benchmark, particularly those made by users in prompts. Complementing Sharma et al. (2023), we construct the prompt by posing from the beginning a poem description and revealing the author's name (deliberately incorrect). We use this strategy to focus on the importance of the task requested at the top of the input ((Sharma et al., 2023) asked for information or expressed opinions at the end). Therefore, we collect ten English poems (Table 7) and 60 authors (Table 6). Hence, we produce 600 prompts using the formula "Describe this {wrong author} poem:". We consider the answer where LLMs solve the task by mentioning the author present in the prompt for the given poem as sycophantic.

¹In additional experiments (Appendix K), we show that evaluation with an LLM as judge delivers comparable results

Evaluation We evaluate the percentage of responses where the model describes the answered poem under the name of the author provided. For example, in Figure 3, all LLMs except GPT-3.5 generate a poem description using the author’s name mentioned in the input. Conversely, GPT-3.5 mentions a different poem from the one requested. In this case, we do not consider his response an error and, consequently, sycophantic behaviour.

3.3 Measuring LLMs Self-Confidence

Finally, to estimate the LLMs’ confidence to deliver correct answers although the user provides misleading hints (§2.3), we use the following benchmarks:

General Commonsense Reasoning: We use CommonSenseQA (CSQA) and OpenBookQA (Mihaylov et al., 2018) (OBQA). CommonSenseQA deals with different types of general commonsense knowledge, while OpenBookQA is a resource that contains questions related to common knowledge and rich text comprehension. High school-level open-book exams inspire it in physics and biology.

Physical Interaction: We use Physical Interaction Question Answering (PIQA) (Bisk et al., 2019), which is a resource consisting of a series of everyday situations with a pair of typical and atypical solutions.

Social Interaction: We use the Social Interaction Question Answering (SIQA) (?) benchmark that is focused on reasoning about people’s actions and social implications. The actions in Social IQA cover various social situations and candidates for plausible and not plausible answers.

Math Word Problem: We select two similar benchmarks: GSM8K (Cobbe et al., 2021), Multi-Arith (Roy and Roth, 2015). In Math Word Problems, there is a textual input, a mathematical problem with a number as its target value. We, therefore, construct the hints by systematically entering the correct and incorrect numerical targets, as shown in Figure 2 in Appendix C. In the case of correct hints, we use real numerical targets; instead, in the case of incorrect hints, we insert a relatively small numerical random bias, as described in Appendix C.

Evaluation To observe the LLMs’ Self-confidence and robustness to misleading interventions, we evaluate the LLMs’ accuracy

(string matching between target and answer) and percentage of agreement with the hint provided by the human in the prompt.

3.4 Models

To analyse the sycophantic behaviours of state-of-the-art fine-tuned LLMs, we experiment with three groups of models: two from the OpenAI family (OpenAI, 2023): GPT-3.5 and GPT-4; two forms of the Meta family (Touvron et al., 2023): Llama2-chat-7b, Llama2-chat-13b, and Llama2-chat-70b; and two forms of the Mistral family (Jiang et al., 2023, 2024): Mistral-7b and Mixtral-8x7b. To simplify the discussion, we omit "chat" and the letter "b" for the Meta and Mistral families. The resulting names will be Llama2-7, -13, -70, Mistral-7, and Mixtral-8x7. We used both open-source models - the Meta and Mistral families - to make our work more reproducible and closed-source models - the OpenAI family - because they demonstrate outstanding performance in many NLP tasks. In Appendix D, we better describe the characteristics of the models and all the parameters adopted for our experiments.

4 Results & Discussion

Large Language Models (LLMs) fine-tuned via human feedback appear sensitive to user prompts. The proposed models, although at different scales, seem to follow user beliefs, even though these were provided in opposite directions on topics of politics and philosophy, as discussed in §4.1. In addition, they easily fall into the trick of mimicking user mistakes, as demonstrated in §4.2.

Finally, although previous experiments argue a variety of weaknesses and a lack of robustness and firmness on the part of LLMs, the sycophantic behaviour towards user interactions appears significantly lower in question-answering tasks. In fact, despite misleading hints, the examined models appear self-confident in their choices, as deeply analysed in §4.3.

4.1 Chameleon LLMs

The human point-of-view manifested via users’ beliefs in prompts tends not to be contrasted by LLMs that reveal a chameleon-like attitude. By systematically proposing a series of prompts on the same topics with different opinions, LLMs generate responses that align with the views expressed by the users, even when these are completely contradic-




Figure 4: We investigate the tendency of LLMs to repeat user opinions (sycophancy). Using three benchmark beliefs (§3.1), we estimate the percentage of model responses in agreement with the users’ point-of-view.

tory. This can be seen in Figure 5 and scaled on different models as reported in Appendix E.

The results in Figure 4 show that for Political Questions (POLI-Q) and NLP Research Questions (NLP-Q) related topics, the difference in agreement between the GPTs and Llama2-70 is about 4 points, and with Llama2-13 and -7 is about 8 points on average. On the other hand, the Mistral models have an average agreement rate of around 62%. In contrast, in Philosophy-related Question (PHI-Q), we observe the gap only in the Mistrals, which reveal an agreement score of around 72%. Similar but larger scaled values are also present in the Llama models, which is not the case in the GPTs. We believe this is a stronger indication that LLMs produce chameleon-like responses when conversing with humans and expressing their points of view.

Moreover, by analysing in-family attitudes, it is possible to observe that LLMs with more parameters reveal higher agreement answers than those with fewer parameters; see Llama2-7,-13, Llama2-70 and both GPTs (in Figure 4) and Appendix E. This phenomenon appears to be directly related to fine-tuning technique, as claimed on a smaller scale with Reinforcement Learning from Human Feedback method in (Perez et al., 2022). We observe the same phenomenon in the Mistral models, which, although using a different refinement technique called Direct Performance Optimisation (Rafailov et al., 2023), are nevertheless trained to maximise human preferences score. Finally, to reduce the biased generations, we repeated the experiments, omitting the part about agreement (i.e. ‘Do you agree?’), as shown in Appendix K. The results in Table 10 are comparable with those obtained in the version of the task introduced in §3.1.

Downstream of the results discussed and confirmed by experiments on additional models (Ap-

pendix E), it can be observed that the sycophantic attitudes exhibited by LLMs refined via the RLHF technique are absorbed more by models with high numbers of parameters.

However, these behaviours are on topics that do not necessarily have an unquestionable target response. Therefore, in §4.2, we analyse the behaviour of the LLMs that perform a task related to a potentially misplaced prompt.




Figure 5: We investigate the agreement rate with user mistakes in our benchmark (§3.2). The considered LLMs tend to mimic human mistakes also when faced with actual error as in Figure 3.

4.2 When LLMs fall into mistakes

Even the more robust LLMs seem not to contradict the users’ point of view by generating answers in agreement with it, as revealed in §4.1. The results may seem positive, as a model should not be biased under specific topics such as politics, but there may also be weaknesses. Indeed, prioritising the human point of view by taking the prompt as correct could easily lead LLMs into error. One possible weakness can be observed in Figure 3 on one instance of the Non-Contradiction benchmark and in larger experiments in Figure 5. It appears that the LLMs primarily focus on performing the task required by the user rather than pointing out errors in the prompt. By systematically asking for a text description written by purposely mistaken authors, we observe the tendency of the LLMs to solve the task without actually addressing the truthfulness of the question (see, for instance, Figure 3).

However, although the general results tend to mimic user errors, the fine-grained results show that the phenomenon is considerably lower when the error is significant, as discussed in Appendix F. In the fine-grained analysis, we discern between especially erroneous prompts with entities related to real poets and several public personages. The